

# OpenMP for Computational Scientists

## 1: Parallel worksharing

Dr Tom Deakin  
University of Bristol

Tuesday 1 December, 2020



# What is OpenMP?

A collection of compiler directives, library routines, and environment variables for parallelism for shared memory parallel programs.

- ▶ Create and manage parallel programs while permitting portability.
- ▶ User-directed parallelization.

A *specification* of annotations you can make to your program in order to make it parallel.

- ▶ OpenMP mostly formed of *compiler directives*

```
!$omp construct [clause [clause]...]
```

These tell the compiler to insert some extra code on your behalf.

- ▶ Compiler directives usually apply to a *structured block* of statements. Limited scoping in Fortran means we often need to use *end* directives.

```
!$omp construct  
... ! lines of Fortran code  
!$omp end construct
```

- ▶ Library API calls

```
use omp_lib  
call omp_...()
```

Turn on OpenMP in the compiler:

```
gfortran *.f90 -fopenmp # GNU  
ifort *.f90 -qopenmp    # Intel  
ftn *.f90 -homp         # Cray (now off by default)  
pgf90 *.f90 -mp         # PGI
```

To also use the API calls within the code, use the library:

```
USE omp_lib
```


## Note

No need to include the library if only using the compiler directives. The library only gets you the API calls.

## Shared memory

OpenMP is for shared memory programming: all threads have access to a shared address space.

A typical HPC node consisting of 2 multi-core CPUs.




*All* threads (each running on a core) can access the same memory.

Different to MPI, where one process cannot see the memory of another without explicit

OpenMP is for shared memory programming: all threads have access to a shared address space.

A typical HPC node consisting of 2 multi-core CPUs.



*All* threads (each running on a core) can access the same memory.

Different to MPI, where one process cannot see the memory of another without explicit 5 / 27


Serial/sequential execution:



Serial/sequential execution:



In a *fork-join* model, code starts serial, *forks* a *team* of threads then *joins* them back to serial execution.



Nested threads are allowed, where a thread forks its own team of threads.



```
1  program hello
2
3  !$omp parallel
4      print *, "Hello"
5  !$omp end parallel
6
7  end program hello
```

Threads *redundantly* execute code in the block.

Each thread will output Hello.

Threads are synchronised at the end of the parallel region.

You might need to set the number of threads to launch (though typically you'll leave OpenMP to set the number of threads for you at run-time).

OpenMP has 3 ways to do this:

- ▶ Environment variables

```
OMP_NUM_THREADS=16
```

- ▶ API calls

```
call omp_set_num_threads(16)
```

- ▶ Clauses

```
!$omp parallel num_threads(16)  
!$omp end parallel
```

In general it's better to use environment variables if you need to do this, as this approach gives you more flexibility at runtime.

Parallel programs often written in a SPMD style:

**Single Program, Multiple Data.**

- ▶ MPI has a SPMD model.
- ▶ Threads run the same code, and use their ID to work out which data to operate on.

The OpenMP API gives you calls to determine thread information when *inside* a parallel region:

- ▶ Get number of threads  
`nthreads = omp_get_num_threads()`
- ▶ Get thread ID  
`tid = omp_get_thread_num()`

## Walkthrough parallelising vector addition using OpenMP.

```
1  program vecadd
2      integer :: N = 1024  ! Length of array
3      ! Arrays
4      real(kind=8), allocatable, dimension(:) :: A, B, C
5      integer :: i          ! Loop counter
6
7      ! Allocate and initialise vectors
8      allocate(A(N), B(N), C(N))
9      A = 1.0; B = 2.0; C = 0.0
10
11     ! Vector add
12     do i = 1, N
13         C(i) = A(i) + B(i)
14     end do
15
16     deallocate(A,B,C)
17 end program vecadd
```

Add parallel region around work

```
!$omp parallel  
do i = 1, N  
    C(i) = A(i) + B(i)  
end do  
!$omp end parallel
```

Every thread will now do the entire vector addition — redundantly!

### Get thread IDs

```
integer :: tid, nthreads

!$omp parallel
tid = omp_get_thread_num()
nthreads = omp_get_num_threads()

do i = 1, N
    C(i) = A(i) + B(i)
end do
!$omp end parallel
```

### Get thread IDs

```
integer :: tid, nthreads

!$omp parallel
tid = omp_get_thread_num()
nthreads = omp_get_num_threads()

do i = 1, N
    C(i) = A(i) + B(i)
end do
!$omp end parallel
```

Incorrect behaviour at runtime

What's the problem here?

- ▶ In OpenMP, all variables are *shared* between threads.
- ▶ But each thread needs its own copy of `tid`.
- ▶ Solution: use the `private` clause on the `parallel` region.
- ▶ This gives each thread its own unique copy in memory for the variable.

```
integer :: tid, nthreads

!$omp parallel private(tid)
tid = omp_get_thread_num()
nthreads = omp_get_num_threads()

do i = 1, N
  C(i) = A(i) + B(i)
end do
!$omp end parallel
```

Much more information about data sharing clauses in next session.



Finally, distribute the iteration space across the threads.

```
integer :: tid, nthreads

!$omp parallel private(tid)
tid = omp_get_thread_num()
nthreads = omp_get_num_threads()

do i = 1+(tid*N/nthreads), (tid+1)*N/nthreads
    C(i) = A(i) + B(i)
end do
!$omp end parallel
```

### Remember

Thread IDs are numbered from 0 in OpenMP. Be careful with your index calculation.

A barrier simply synchronises threads in a parallel region.

```
1  !$omp parallel private(tid)
2
3  tid = omp_get_thread_num()
4  A(tid) = big_work1(tid)
5
6  !$omp barrier
7
8  B(tid) = big_work2(A, tid)
9
10 !$omp end parallel
```

- ▶ Running in parallel, need to compute A(:) before computing B(:).
- ▶ The barrier ensures all threads wait between these statements.
- ▶ Must ensure all threads encounter the barrier.

- ▶ The SPMD approach requires lots of bookkeeping.
- ▶ Common pattern of splitting loop iterations between threads.
- ▶ OpenMP has worksharing constructs to help with this.
- ▶ Used within a parallel region.
- ▶ The loop iterator is made **private** by default: no need for data sharing clause.

```
!$omp parallel
!$omp do
do i = 1, N
  C(i) = A(i) + B(i)
end do
!$omp end do
!$omp end parallel
```

Implicit synchronisation point at the *!\$omp end do*.

Generally it's convenient to combine the directives:

```
!$omp parallel do  
do i = 1, N  
    ... ! loop body  
end do  
!$omp end parallel do
```

- ▶ This starts a parallel region, forking some threads.
- ▶ Each thread then gets a portion of the iteration space and computes the loop body in parallel.
- ▶ Implicit synchronisation point at the `end do`.
- ▶ Threads finally join again; later code executes sequentially.

The vector add codes are available in the repository for you to look at:

- ▶ Serial: `vadd.f90`
- ▶ SPMD: `vadd_spmd.f90`
- ▶ Worksharing: `vadd_parallel_eldo.f90`


- ▶ Often have tightly nested loops in your code.
- ▶ E.g. 2D grid code, every cell is independent.
- ▶ OpenMP worksharing would only parallelise over first loop with each thread performing inner loop serially.
- ▶ Use the `collapse(...)` clause to combine iteration spaces.
- ▶ OpenMP then workshares the combined iteration space.

```
!$omp parallel do collapse(2)  
do i = 1, N  
  do j = 1, N  
    ... ! loop body  
  end do  
end do  
!$omp end parallel do
```

All  $N^2$  iterations are distributed across threads, rather than just the  $N$  of the outer loop.

## 5-point stencil exercise

First exercise: parallelise a simple 5-point stencil code using OpenMP.



Value in every cell is set to the average of its neighbours.

Take `stencil.f90` and parallelise it using OpenMP:

1. Using a SPMD style.
2. Using the OpenMP worksharing clauses.
3. Vary the number of threads using `OMP_NUM_THREADS`.

Focus on parallelising the main loop(s):

```
do i = 1, nx
  do j = 1, ny
    Anew(i,j) = (A(i-1,j) + A(i+1,j) + A(i,j) + A(i,j-1) + A(i,j+1))
    ↪ / 5.0
  end do
end do
```

Sample solutions are provided, but do try it yourself first.



- ▶ The worksharing clauses use default rules for assigning iterations to threads.
- ▶ Can use the schedule clause to specify the distribution.
- ▶ General format:

```
!$omp parallel do schedule(...)
```

Next slides go through the options, using the following loop as an example:

```
!$omp parallel do num_threads(4)  
do i = 1, 100  
    ... ! loop body  
end do  
!$omp end parallel do
```

```
schedule(static)  
schedule(static,16)
```

- ▶ Static schedule divides iterations into chunks and assigns chunks to threads in round-robin.
- ▶ If no chunk size specified, iteration space divided roughly equally.

For our example loop:  
`schedule(static)`

| Thread ID | Iterations |
|-----------|------------|
| 0         | 1–25       |
| 1         | 26–50      |
| 2         | 51–75      |
| 3         | 76–100     |

`schedule(static,16)`

| Thread ID | Iterations    |
|-----------|---------------|
| 0         | 1–16, 65–80   |
| 1         | 17–32, 81–96  |
| 2         | 33–48, 97–100 |
| 3         | 49–64         |

```
schedule(dynamic)
```

```
schedule(dynamic,16)
```

- ▶ Iteration space is divided into chunks according to chunk size.
- ▶ If no chunk size specified, default size is one.
- ▶ Each thread requests and executes a chunk, until no more chunks remain.
- ▶ Useful for unbalanced work-loads if some threads complete work faster.

For our example with a chunk size of 16:

- ▶ The iteration space is split into chunk of 16 (the last chunk may be smaller).
- ▶ Each thread gets one chunk, then requests a new chunk to work on.

```
schedule(guided)  
schedule(guided,16)
```

- ▶ Similar to assignment to `dynamic`, except the chunk size decreases over time.
- ▶ Granularity of work chunks gets finer over time.
- ▶ If no chunk size is specified, the default size is one.
- ▶ Useful to try to mitigate overheads of a `dynamic` schedule by starting with large chunks of work.

For our example with a chunk size of 16:

- ▶ Each thread gets a chunk of 16 to work on.
- ▶ Each thread requests a new chunk, which might be smaller than 16.

`schedule(auto)`

- ▶ Let the compiler or runtime choose the schedule.

`schedule(runtime)`

- ▶ Get the schedule from the `OMP_SCHEDULE` environment variable.

### Recommendation

Just use a static schedule unless there is a good reason not to! `static` is usually the fastest of all the options. The choice of schedules is an advanced tuning option.

- ▶ May have series of loops in your code which are independent.
- ▶ Threads must wait/synchronise at the end of the loop.
- ▶ But it might be possible to delay this synchronisation using the `nowait` clause.
- ▶ When a thread finishes the first loop, it starts on the next loop.

```
1  !$omp parallel
2  !$omp do nowait
3  do i = 1, N
4      A(i) = i
5  end do
6  !$omp end do          ! No barrier!
7  !$omp do
8  do i = 1, N
9      B(i) = i
10 end do
11 !$omp end do          ! Implicit barrier
12 !$omp end parallel ! Implicit barrier
```