

OpenMP for Computational Scientists

6: Tasking and Tools

Tom Deakin
University of Bristol

- ▶ Quick recap of **target** exercise
- ▶ Poor man's tasking: sections
- ▶ The single and master constructs
- ▶ Tasking in OpenMP
 - ▶ Generating tasks
 - ▶ Synchronising tasks
 - ▶ Data sharing rules
 - ▶ Specifying task dependencies
- ▶ Tools
- ▶ OpenMP 5.0

Take your 5-point stencil code and run on a GPU with the **target** directive.

1. Copy data to/from the device.

```
!$omp target enter data map(to: A, Atmp)

do t = 1, ntimes

    call stencil(...)

end do

!$omp target exit data map(from: A, Atmp)
```

2. Offload parallel execution of the loops.

```
!$omp target map(tofrom:total)
!$omp teams distribute parallel do reduction(+:total) collapse(2)
do j = 1, ny
  do i = 1, nx
    Atmp(i,j) = (A(i-1,j) + A(i+1,j) + A(i,j) + A(i,j-1) + A(i,j+1)) *
      ↪ 0.2
    total = total + Atmp(i,j)
  end do
end do
!$omp end teams distribute parallel do
!$omp end target
```

Don't forget to copy back the reduction variable.

- ▶ What if your code doesn't follow a standard parallel loop pattern?
- ▶ OpenMP needs to know the loop count at runtime.
- ▶ Recursive and tree/graph based algorithms inconvenient to program with parallel loops.
- ▶ Tasking allows you to package work and data into units (tasks) and have them scheduled in parallel.

Example: processing a linked list:

```
p => head
do while(associated(p))
  call process(p)
  p => p%next
end do
```

When not to use tasks

Think!

If you have a regular parallel algorithm, tasks are *unlikely* to be the best approach.

Use the worksharing directives instead.

- ▶ Sections give you a way to assign different work to different threads.
- ▶ Useful for a producer/consumer pattern, but not really recursive algorithms.
- ▶ If fewer sections than threads, threads sit idle.
- ▶ If more sections than threads, sections assigned by implementation.

```
!$omp parallel sections  
    !$omp section  
    call work1()  
  
    !$omp section  
    call work2()  
  
!$omp end parallel sections
```

- ▶ Often necessary for only one thread to do some work in a parallel region.
- ▶ Wrapping code with the `single` constructs means only one thread in the parallel region will execute that code.
- ▶ Not defined which thread will actually execute the block.
- ▶ There is an implicit barrier for this construct (can avoid with `nowait` clause).

```
!$omp parallel  
  call do_work()  
  !$omp single  
  call exchange_halos()  
  !$omp end single  
!$omp end parallel
```


- ▶ Similar to single construct, except guaranteed the master thread will execute the block.
- ▶ There is *no* implicit barrier for this construct!
- ▶ Useful for cases where synchronisation isn't required:
 - ▶ Printing out messages to the screen.
 - ▶ **Generating tasks!**

```
!$omp parallel
  call do_work()
  if (conv .eq. .true.)
    !$omp master
    print *, "Converged!"
    !$omp end master
  exit
end if
!$omp end parallel
```


- ▶ OpenMP 3.0 introduced real tasking.
- ▶ Later versions refined and added features.

```
!$omp parallel
  !$omp master
    !$omp task
    call do_task(x,y,z)
    !$omp end task

    !$omp task
    call do_task(i,j,k)
    !$omp end task
  !$omp end master
!$omp end parallel
```

- ▶ Use master thread to generate tasks for all threads to do.
- ▶ Idle threads take tasks off the task queue.
- ▶ Implicit barrier at the **end parallel** ensures all tasks complete.

Useful to think about tasks organised as a tree (or graph).



- ▶ The Initial Task is the *parent* of Task A.
- ▶ Task A is the *child* of the Initial Task.
- ▶ Task A and Task B are *siblings*.

```
!$omp parallel
  !$omp master
    !$omp task
    call do_task(x,y,z)
    !$omp end task

    !$omp task
    call do_task(i,j,k)
    !$omp end task
  !$omp end master
!$omp end parallel
```

```
!$omp parallel
  !$omp master
    !$omp task
      call do_task(x,y,z)
    !$omp end task


    !$omp task
      call do_task(i,j,k)
    !$omp end task
  !$omp end master
!$omp end parallel
```

Initial Task

1. Initial thread begins serial execution.

```
!$omp parallel
  !$omp master
    !$omp task
      call do_task(x,y,z)
    !$omp end task


    !$omp task
      call do_task(i,j,k)
    !$omp end task
  !$omp end master
!$omp end parallel
```



1. Initial thread begins serial execution.
2. Parallel region encountered, implicit task generated per thread.

```
!$omp parallel
  !$omp master
    !$omp task
      call do_task(x,y,z)
    !$omp end task


    !$omp task
      call do_task(i,j,k)
    !$omp end task
  !$omp end master
!$omp end parallel
```



1. Initial thread begins serial execution.
2. Parallel region encountered, implicit task generated per thread.
3. Master thread T0 generates tasks. Tasks are *children* of T0, and *siblings* of each other.

```
!$omp parallel
  !$omp master
    !$omp task
      call do_task(x,y,z)
    !$omp end task


    !$omp task
      call do_task(i,j,k)
    !$omp end task
  !$omp end master
!$omp end parallel
```



1. Initial thread begins serial execution.
2. Parallel region encountered, implicit task generated per thread.
3. Master thread T0 generates tasks. Tasks are *children* of T0, and *siblings* of each other.
4. When Tasks T0 and T1 are done they reach the implicit barrier. Threads free to execute other tasks.



```
!$omp parallel
  !$omp master
    !$omp task
      call do_task(x,y,z)
    !$omp end task

    !$omp task
      call do_task(i,j,k)
    !$omp end task
  !$omp end master
!$omp end parallel
```



1. Initial thread begins serial execution.
2. Parallel region encountered, implicit task generated per thread.
3. Master thread T0 generates tasks. Tasks are *children* of T0, and *siblings* of each other.
4. When Tasks T0 and T1 are done they reach the implicit barrier. Threads free to execute other tasks.
5. Initial task resumes after implicit barrier at **end parallel**.

- ▶ It's useful to think about how the OpenMP runtime itself implements tasking.
- ▶ In general, the runtime maintains a queue of tasks.
- ▶ Threads enqueue and dequeue tasks from the queue.



- ▶ Intel/Clang runtime has one task queue per thread, and allows work-stealing.
- ▶ Gives lower contention than case where all threads access a single queue.

Three places where tasks synchronise:

1. At thread barriers:

- ▶ Implicit barriers (like at end of `parallel` and `single` regions).
- ▶ Explicit barrier constructs.
- ▶ *All* previously generated tasks will complete.

2. At the `taskwait` construct:


- ▶ Waits on all child tasks of the current task before continuing.
- ▶ Only applies to tasks that the current task *generated*.
- ▶ Everything in OpenMP is defined as task, so here the “outer/current” task is the one that called the first task construct.

3. At the end of the `taskgroup` construct:

- ▶ Waits on all child tasks generated within the `taskgroup` region.
- ▶ *Includes* waiting on descendants.
- ▶ Like `taskwait`, but allows grouping of child tasks.

!\$omp taskwait: waits for *generated* (child) tasks to finish.

```
subroutine do_a()  
  !$omp task  
  call do_b()  
  !$omp end task  
  
  !$omp task  
  call do_c()  
  !$omp end task  
end subroutine  
  
subroutine do_b()  
  !$omp task  
  call do_d()  
  !$omp end task  
  
  !$omp task  
  call do_e()  
  !$omp end task  
end subroutine
```



`!$omp taskwait`: waits for *generated* (child) tasks to finish.

```
subroutine do_a()  
  !$omp task  
  call do_b()  
  !$omp end task  
  
  !$omp task  
  call do_c()  
  !$omp end task  
end subroutine  
  
subroutine do_b()  
  !$omp task  
  call do_d()  
  !$omp end task  
  
  !$omp task  
  call do_e()  
  !$omp end task  
end subroutine
```

```
$ OMP_NUM_THREADS=4 ./tasks  
Task A starting  
Task A finished  
Task C starting  
Task B starting  
Task C finished  
Task B finished  
Task D starting  
Task E starting  
Task E finished  
Task D finished
```

Tasks finish even when child tasks haven't started.

- ▶ A `taskwait` in Task A will wait for Tasks B and C to finish.
- ▶ Task B finishes when Tasks D and E have been *generated*, not completed.
- ▶ So this `taskwait` *will not* wait for Tasks D and E to finish.

```
subroutine do_a()  
  !$omp task  
  call do_b()  
  !$omp end task  
  
  !$omp task  
  call do_c()  
  !$omp end task  
  
  !$omp taskwait  
end subroutine
```

```
$ OMP_NUM_THREADS=4 ./tasks  
Task A starting  
Task B starting  
Task C starting  
Task B finished  
Task C finished  
Task D starting  
Task A finished  
Task E starting  
Task D finished  
Task E finished
```



- ▶ A `taskwait` *also* in Task B will ensure Task B waits for Tasks D and E to finish before finishing itself.
- ▶ Task A waits for Tasks B to finish, which is waiting for Tasks D and E to finish.
- ▶ So Task A will now finish when Tasks B, C, D and E are finished.

```
subroutine do_b()  
  !$omp task  
  call do_d()  
  !$omp end task  
  
  !$omp task  
  call do_e()  
  !$omp end task  
  
  !$omp taskwait  
end subroutine
```

```
$ OMP_NUM_THREADS=4 ./tasks  
Task A starting  
Task B starting  
Task C starting  
Task E starting  
Task D starting  
Task C finished  
Task E finished  
Task D finished  
Task B finished  
Task A finished
```

- ▶ Similar to `taskwait` but also waits on children of child tasks.
- ▶ Creates a set of tasks containing *all* generated tasks, including descendants, in the region.
- ▶ Encountering task suspended until all the tasks in the set complete.

```
!$omp taskgroup
```

```
! Generate some tasks...
```

```
!$omp end taskgroup
```

```
! Here, all tasks generated in the taskgroup region  
! have finished
```



```
1  !$omp parallel
2  !$omp single
3
4  !$omp task
5  call background_work()
6  !$omp end task
7
8  !$omp taskgroup
9      !$omp task
10     call recursive_work()
11     !$omp end task
12 !$omp end taskgroup
13
14 call check_work()
15
16 !$omp end single
17 !$omp end parallel
```

- ▶ Using single region to generate tasks, which has implicit barrier.
- ▶ Start a background task, and perform a recursive task which generates more tasks.
- ▶ The taskgroup waits for all the tasks in the group only, *excluding* the background task.
- ▶ Background task can continue while executing `check_work()`.
- ▶ If used a `taskwait` before line 14, it would also wait for the *background* task. Also need to ensure wait on recursive task's descendants.

- ▶ Can use these three data sharing clauses from `parallel` regions on the `task` construct.
- ▶ The definitions are really the same, but applied to tasks, not threads.
- ▶ Need to think carefully about the appropriate clauses for each variable.
- ▶ `shared(x)`: There is one copy of the `x` variable, shared between the current and child tasks.
- ▶ `private(x)`: Each task gets its own local `x` variable. It is not initialised.
- ▶ `firstprivate(x)`: Each task gets its own local `x` variable. It is initialised to the value of the original variable when *encountered*.

- ▶ Tasks and parallel regions have different default data sharing rules.
- ▶ By default, data is *shared* for `parallel` regions.
- ▶ By default, data is *firstprivate* for task constructs.
 - ▶ Unless, it's shared by the enclosing (the outer) region.
- ▶ Sensible behaviour because task execution may be delayed, and original variables may no longer be in scope.
- ▶ Using `default(none)` is recommended here especially.

```
!$omp parallel shared(A) private(B)  
  
    !$omp task  
        call compute(A,B)  
    !$omp end task  
  
!$omp end parallel
```

Within the task:

- ▶ A is shared.
- ▶ B is firstprivate.
- ▶ Any variables inside the subroutine are **private**.

```
1  recursive integer function fib(n) result(res)
2      integer :: n, i, j
3      if (n .lt. 2) then
4          res = n
5      else
6          i = fib(n-1)
7          j = fib(n-2)
8
9          res = i+j
10     end if
11 end function
12
13 program fibonacci
14     integer :: res
15
16     res = fib(40)
17     print *, res
18 end program
```

Note: there are better ways to calculate Fibonacci numbers...

Create the parallel region:

```
1  program fibonacci
2
3      integer :: res
4
5      !$omp parallel
6          !$omp master
7              res = fib(40)
8          !$omp end master
9      !$omp end parallel
10
11     print *, res
12
13 end program
```


```
recursive integer function
↪  fib(n) result(res)
    integer :: n, i, j
    if (n .lt. 2) then
        res = n
    else
        !$omp task shared(i)
        i = fib(n-1)
        !$omp end task
        !$omp task shared(j)
        j = fib(n-2)
        !$omp end task
        !$omp taskwait
        res = i+j
    end if
end function
```

- ▶ Recursively creates a binary tree of tasks.
- ▶ This first task `fib(40)` creates two tasks: `fib(39)` and `fib(38)`.
- ▶ The `taskwait` construct ensures that the child tasks are finished before summing their return values.
- ▶ `i` and `j` *must* be shared so that the results are retained by the calling task.
- ▶ They are local variables so would have been `firstprivate` for the tasks by default.

- ▶ OpenMP gives you the ability to define an ordering of tasks based on input and output data dependencies.
- ▶ In other words, can make tasks wait for other sibling tasks to finish before starting.
- ▶ Need to specify input and output dependencies.
- ▶ Only checks for previously generated siblings with a dependency specified.
- ▶ Add a `depend` clause to the task directive:
 - ▶ `depend(in: A)`: task depends on siblings with `out` or `inout` dependency `A`.
 - ▶ `depend(out: A)`: task depends on siblings with any dependency on `A`.
 - ▶ `depend(inout: A)`: same as `depend(out: A)`.

Dependency example

```
1  !$omp task depend(out: x)
2  x = 1.0
3  !$omp end task
4
5  !$omp task depend(in: x) depend(out: y)
6  y = f(x)
7  !$omp end task
8
9  !$omp task depend(in: x) depend(out: z)
10 z = g(x)
11 !$omp end task
12
13 !$omp task depend(in: y, z)
14 print *, y+z
15 !$omp end task
```



Dependency example

```
1  !$omp task depend(out: x)
2  x = 1.0
3  !$omp end task
4
5  !$omp task depend(in: x)
6  ↪ depend(out: y)
7  y = f(x)
8  !$omp end task
9
10 !$omp task depend(in: x)
11 ↪ depend(out: z)
12 z = g(x)
13 !$omp end task
14
15 !$omp task depend(in: y, z)
16 print *, y+z
17 !$omp end task
```

- ▶ Must specify first out dependency.
- ▶ Then generate $y = f(x)$ tasks which depends on it.
- ▶ Then generate $z = g(x)$ task, which depends on first but not second.
- ▶ Generate task which depends on middle two tasks.
- ▶ Note, OpenMP still sees all these tasks as *siblings*.

OpenMP 5.0 will introduce the `mutexinoutset` dependency qualifier.

- ▶ Ensures commutativity between tasks depending on `x`.
- ▶ Sibling tasks with `depend(inout: x)` would need to be run in the order the tasks were generated.
- ▶ `mutexinoutset` means they can run in any order, not just the generated order.
- ▶ Tasks still won't run in parallel.

```
!$omp task (inout: y)
call slow(y) ! Task A
!$omp end task
```

```
!$omp task (inout: z)
call fast(z) ! Task B
!$omp end task
```

```
!$omp task depend(inout:
  ↪ res) depend(in: y)
res = res + y ! Task C
!$omp end task
```

```
!$omp task depend(inout:
  ↪ res) depend(in: z)
res = res + z ! Task D
!$omp end task
```

- ▶ Task A and B will start.
- ▶ Task C can start once Task A has finished.
- ▶ Task D can start once both Tasks B and C have finished.
- ▶ Ideally should just wait on Task B, not C.
- ▶ Solved by replacing `depend(inout: res)` with `depend(mutexinoutset: res)`.

- ▶ Once a task starts executing, it goes to completion unless it encounters a *task scheduling point*.
- ▶ At these points, the threads can pick another task to start executing or resume a task.
- ▶ These points occur when you:
 - ▶ Generate a task
 - ▶ Finish a task
 - ▶ Reach the end of a taskgroup construct.
 - ▶ Synchronise with barrier or taskwait.
 - ▶ And, at a taskyield construct.
- ▶ The taskyield construct allows tasks to be suspended so another task can start.
- ▶ Really useful when polling for a lock, otherwise the task never “sleeps”.
- ▶ Not all implementations actually suspend tasks (Intel/Clang does).

```
1  !$omp task
2    call do_something()
3
4    ! Wait for the lock
5    do while (.not. omp_test_lock(lock))
6      ! Couldn't get lock, so suspend task and schedule
       ↪ another
7      !$omp taskyield
8    end do
9
10   ! Only one task should call this function
11   call critical_code()
12
13   ! Release the lock
14   call omp_unset_lock(lock)
15 !$omp end task
```

- ▶ Generates a task for each loop iteration.
- ▶ Shorthand for generating tasks in a tiled `parallel do` loop.
- ▶ Tasks form part of a `taskgroup`, unless use the `nogroup` clause.
- ▶ `grainsize(num)` clause: sets the minimum number of iterations for each generated task.
- ▶ `num_tasks(num)` clause: set number of tasks to create. Unlikely to want to specify this in practice.
- ▶ `collapse` clause can also be used.

```
!$omp taskloop  
do i = 1, N  
    call do_work(i)  
end do  
!$omp end taskloop
```

Some rare clauses you can use when generating tasks, but unlikely to use them in most cases.

- ▶ `untied`: If the task gets suspended, any thread can start it. Without this clause, the thread which started the task must resume it. Helpful for load balancing when used with `taskyield`.
- ▶ `mergeable`: The task can be merged into the parent task, sharing its data environment. Execution might happen straight away, or be delayed.
- ▶ `priority(n)`: The task is given a priority value $n \geq 0$. Gives a hint to the runtime about which tasks to run first.
- ▶ `final(expr)`: If the expression is true, this and all further child tasks are included in the parent task, and just executed immediately. Useful for specifying a minimum tasks size in recursive algorithms.

- ▶ Getting the correct data sharing clauses can be tricky.
- ▶ Don't use tasks for patterns supported by other parts of OpenMP (parallel loops).
- ▶ Tasking comes with overheads.
- ▶ The runtime is good, but can't work miracles.
- ▶ Best results obtained where programmer controls number and granularity of tasks.

Intel Parallel Studio XE

- ▶ Intel Advisor
 - ▶ Shows vectorisation efficiency and Roofline analysis.
 - ▶ Can show you where might be good to thread, but obviously can't refactor your code for you to make it parallelisable.
- ▶ Intel Inspector
 - ▶ Debugger for memory errors and thread errors (such as deadlock).
- ▶ Intel vTune Amplifier
 - ▶ Performance tuning using hardware counters and metrics.
 - ▶ Very detailed information presented, but tries to summarise it.
- ▶ The documentation on the Intel website is really helpful.

ARM HPC Tools

- ▶ From ARM, but cross-platform. Used to be the Allinea Forge tools.
- ▶ DDT: Graphical debugger. Really good for parallel programs in MPI and OpenMP (and both).
- ▶ Map: Simple to use profiler. Can show thread activity, synchronisation and communication overheads.
- ▶ <https://www.arm.com/products/development-tools/hpc-tools/cross-platform/forge>

Cray PerfTools

- ▶ CrayPat
 - ▶ Performance analysis tool for Cray systems.
 - ▶ Gives overview of expensive routines and communication and synchronisation costs.
 - ▶ Also uses hardware counters for more details information.
- ▶ Reveal
 - ▶ Analyses loop dependencies and uses performance data to automatically try to insert OpenMP constructs.
 - ▶ Doesn't do a complete job, still need to check the constructs are correct.
 - ▶ Obviously can't refactor a loop to make is parallelisable.
 - ▶ <http://www.nersc.gov/users/software/performance-and-debugging-tools/craypat/reveal/>

- ▶ Can use the NVIDIA debuggers and profilers for OpenMP `target` constructs.
- ▶ Extrae and Paraver profiler and visualiser from Barcelona Supercomputing Center: <https://tools.bsc.es>.
- ▶ Score-P and Vampir: good for generating timelines of multi-threaded/process code:
<https://www.alcf.anl.gov/files/Vampir.pdf>
- ▶ TAU profiler and program tracer:
<https://www.cs.uoregon.edu/research/tau/tau.ppt>
- ▶ TotalView debugger:
<https://computing.llnl.gov/tutorials/totalview/>

- ▶ OpenMP website: <https://www.openmp.org>
- ▶ cOMPunity: <http://www.compunity.org>
- ▶ Tim Mattson's YouTube tutorial: <https://youtu.be/nE-xN4Bf8XI>
- ▶ SC'08 tutorial from Tim Mattson and Larry Meadows:
<https://openmp.org/mp-documents/omp-hands-on-SC08.pdf>
- ▶ From Lawrence Livermore National Lab:
<https://computing.llnl.gov/tutorials/openMP/>
- ▶ SC'16 Tutorial from Tim Mattson and Alice Koniges:
<https://www.nersc.gov/assets/Uploads/SC16-Programming-Irregular-Applications-with-OpenMP.pdf>
- ▶ OpenMPCon'15 Tutorial from Christian Terboven and Michael Klemm:
<https://openmpcon.org/wp-content/uploads/openmpcon2015-michael-klemm-tasking.pdf>
- ▶ Patrick Atkinson and Simon McIntosh-Smith, *On the Performance of Parallel Tasking Runtimes for an Irregular Fast Multipole Method Application*, IWOMP 2017. https://link.springer.com/chapter/10.1007/978-3-319-65578-9_7.

Due to be released the week before SC'18.

Highlighted additions:

- ▶ `mutexinout` dependency.
- ▶ Task reductions.
- ▶ Scan reduction operations.
- ▶ Task affinity hints to say where tasks might execute based on data.
- ▶ `metadirective` and `declare variant` for selection of directive and function variants based on context.
- ▶ Memory space allocators and traits to support different kinds of memories.
- ▶ `loop construct` for compiler optimisation and parallelisation of loops; specifies iterations can execute in any order.

Plus many others...

Asynchronous offload:

- ▶ Start with the `target` version of the 5-point stencil.
- ▶ Use the `nowait` clause to asynchronously offload part of the computation.
- ▶ Complete the remainder of the work on the host in parallel concurrently with the device.
- ▶ Ensure correct synchronisation with the tasking constructs.

- ▶ Covered a lot of content:
 - ▶ Shared memory parallelism with OpenMP.
 - ▶ Parallelising loops with OpenMP worksharing.
 - ▶ Code optimisations.
 - ▶ Performance analysis, including the Roofline model.
 - ▶ Vectorisation.
 - ▶ NUMA aware programming.
 - ▶ Hybrid MPI+OpenMP codes.
 - ▶ Programming a GPU with OpenMP.
 - ▶ Task based parallelism with OpenMP.
- ▶ Trying the concepts out yourself is the best way of learning a parallel programming model.
- ▶ OpenMP gives you the tools to write performance portable programs.
- ▶ The OpenMP standard, runtimes and compilers continue to develop.